

Game Play Analytics with SQL and Pandas: A Hands-On Guide Using LeetCode 511, 512, 534, 550

Contents

Contents

Chapter 1 – First Login Date (LeetCode 511: Game Play Analysis I)

- 1.1 Problem restatement and data model
- 1.2 Core SQL solution (GROUP BY + MIN)
- 1.3 Core Pandas solution (groupby + min)
- 1.4 Alternative approaches (window functions, subqueries)
- 1.5 Edge cases and data quality considerations
- 1.6 Time and space complexity (SQL engine vs. Pandas)
- 1.7 Code review notes and optimization for large datasets

Chapter 2 – First Login Device (LeetCode 512: Game Play Analysis II)

- 2.1 Problem restatement and subtle points (ties, PK assumption)
- 2.2 Core SQL solution (aggregate + join)
- 2.3 Core Pandas solution (groupby + merge / sort + drop_duplicates)
- 2.4 Alternative window-function solution (ROW_NUMBER)
- 2.5 Edge cases (multiple devices, duplicate rows)
- 2.6 Time and space complexity and indexing
- 2.7 Code review and optimization strategies

Chapter 3 – Cumulative Games So Far (LeetCode 534: Game Play Analysis III)

- 3.1 Problem restatement and “running total” concept
- 3.2 Core SQL solution (window SUM OVER ORDER BY)
- 3.3 Core Pandas solution (sort_values + groupby.cumsum)
- 3.4 Alternative formulations (self-join prefix-sums, correlated subqueries)
- 3.5 Edge cases (missing days, zero games, sparse logins)
- 3.6 Complexity and memory behavior on large tables
- 3.7 Code review: stability, readability, and performance

Chapter 4 – Next-Day Retention After First Login (LeetCode 550: Game Play Analysis IV)

- 4.1 Problem restatement and retention metric definition
- 4.2 Core SQL solution (first login + self-join + COUNT/COUNT)
- 4.3 Core Pandas solution (groupby.min + merge + date arithmetic)
- 4.4 Alternative SQL approaches (EXISTS, DISTINCT, window funcs)
- 4.5 Edge cases (single-day players, time gaps, calendar boundaries)
- 4.6 Complexity and numeric precision (rounding, types)
- 4.7 Code review: correctness, precision, and scalability

Chapter 1 — First Login Date (LeetCode 511: Game Play Analysis I)

1.1 Problem Restatement

We are given a table `Activity`:

Column Name Type

player_id int
device_id int
event_date date
games_played int

(`player_id`, `event_date`) together form the primary key.

Each row represents a login session.

We must determine the first login date (earliest `event_date`) for every player.

Goal:

Return a table with the following structure:

player_id first_login

int date

1.2 SQL Solution (Canonical)

```
SELECT
player_id,
MIN(event_date) AS first_login
FROM Activity
GROUP BY player_id;
```

Explanation

GROUP BY player_id partitions the table by each player.

MIN(event_date) extracts the earliest login date within the group.

Simple and efficient since event_date is comparable and primary key ensures uniqueness.

1.3 Pandas Solution (Canonical)

```
import pandas as pd
```

```
def game_play_analysis_I(activity: pd.DataFrame) -> pd.DataFrame:
result = (
activity
.groupby("player_id", as_index=False)[ "event_date"]
.min()
.rename(columns={ "event_date": "first_login"})
)
return result
```

Explanation (line-by-line)

groupby("player_id") → partitions the DataFrame by player.

["event_date"].min() → computes earliest login date.

rename() ensures the result column matches expected output format.

1.4 Alternative SQL Approaches

Approach A — Window Function + DISTINCT

```
SELECT DISTINCT
player_id,
FIRST_VALUE(event_date) OVER (
PARTITION BY player_id ORDER BY event_date
) AS first_login
FROM Activity;
```

Approach B — Correlated Subquery

```
SELECT
a.player_id,
(SELECT MIN(event_date)
FROM Activity AS b
WHERE a.player_id = b.player_id) AS first_login
```

```
FROM Activity AS a
GROUP BY a.player_id;
```

Approach C — Self-Join on Minimum Date

```
SELECT
a.player_id,
a.event_date AS first_login
FROM Activity AS a
JOIN (
  SELECT player_id, MIN(event_date) AS min_date
  FROM Activity
  GROUP BY player_id
) AS x
ON a.player_id = x.player_id
AND a.event_date = x.min_date;
```

Canonical solution is still the fastest and most readable.

1.5 Alternative Pandas Approaches

Approach A — sort + drop_duplicates

```
def game_play_analysis_l_alt1(activity):
df = activity.sort_values(["player_id", "event_date"])
return df.drop_duplicates("player_id")[["player_id", "event_date"]].rename(
columns={"event_date": "first_login"})
)
```

Approach B — idxmin

```
def game_play_analysis_l_alt2(activity):
idx = activity.groupby("player_id")["event_date"].idxmin()
return activity.loc[idx, ["player_id", "event_date"]].rename(
columns={"event_date": "first_login"})
)
```

1.6 Edge Cases & Discussion

1.6.1 A player logs in once

Both SQL and Pandas handle this naturally—MIN(event_date) is that single date.

1.6.2 Multiple sessions on same date

Not possible due to PK constraint (player_id, event_date).

1.6.3 Extremely large tables

SQL engines use indexes and query planners → generally $O(n \log n)$.

Pandas loads entire dataset into memory → potential memory bottlenecks.

1.6.4 NULL values

Assumed not present (LeetCode guarantees valid input).

If NULL dates existed, MIN would ignore NULLs unless configured otherwise.

1.7 Time & Space Complexity Analysis

SQL Query Complexity

Operation Complexity

Grouping on player_id $O(N \log N)$ typically

Aggregation (MIN) $O(1)$ per group

Total O(N log N)

If the DB index exists on (player_id, event_date), the query may become nearly O(N).

Pandas Complexity

groupby + min

Stage Complexity

Hash-based groupby O(N)

Min aggregation O(1) per group

Total O(N) average

Memory usage:

Pandas must hold the entire dataset in RAM → cost ≈ size of Activity table.

1.8 Code Review Commentary (Optimization Perspective)

SQL Code Review Notes

- ✓ Simple, readable, minimal operations.
- ✓ Uses built-in aggregate MIN → optimal.
- Add an index on (player_id, event_date) for very large datasets.
- ✓ Avoid window functions because unnecessary for this problem.

Pandas Code Review Notes

- ✓ Uses groupby + min (optimal approach).
- ✓ Avoids unnecessary sorting or merging.
- Ensure event_date column is converted to datetime type:
`activity["event_date"] = pd.to_datetime(activity["event_date"])`

Scalability Considerations

For massive datasets (> 50M rows), SQL is preferable due to on-disk processing.

Pandas solution could be rewritten with:

Polars for lazy execution

Dask for distributed processing

1.9 Example Walkthrough

Input

player_id | event_date

1		2016-03-01
1		2016-05-02
2		2017-06-25
3		2016-03-02
3		2018-07-03

Processing

Player 1 → MIN = 2016-03-01

Player 2 → MIN = 2017-06-25

Player 3 → MIN = 2016-03-02

Output

player_id | first_login

1		2016-03-01
2		2017-06-25
3		2016-03-02

End of Chapter 1

```
# Chapter 2 — First Login Device
**LeetCode 512: Game Play Analysis II**
```

2.1 Problem Restatement

We are given the following table:

```
```text
+-----+-----+
| Column Name | Type |
+-----+-----+
| player_id | int |
| device_id | int |
| event_date | date |
| games_played | int |
+-----+-----+
PRIMARY KEY: (player_id, event_date)
```

Each row corresponds to a login session.

Our goal:

For each player, report the device\_id that was used on their FIRST login date.

Expected output:

```
+-----+-----+
| player_id | device_id |
+-----+-----+
```

Because (player\_id, event\_date) is a primary key, there will always be exactly one device\_id corresponding to the earliest date.

## 2.2 SQL Solution (Canonical)

The natural SQL solution uses an aggregate to find the earliest date per player and then joins back to get the device:

```
SELECT
a.player_id,
a.device_id
FROM Activity AS a
JOIN (
SELECT
player_id,
MIN(event_date) AS first_login
FROM Activity
GROUP BY player_id
) AS f
ON a.player_id = f.player_id
AND a.event_date = f.first_login;
```

Why this works

The inner query extracts the earliest login date.

Joining on (player\_id, event\_date) correctly matches the unique login row.

Guaranteed 1-to-1 mapping due to primary key constraint.

## 2.3 Pandas Solution (Canonical)

```

import pandas as pd

def game_play_analysis_II(activity: pd.DataFrame) -> pd.DataFrame:
 first_dates = (
 activity
 .groupby("player_id", as_index=False)[["event_date"]]
 .min()
 .rename(columns={"event_date": "first_login"})
)

 merged = first_dates.merge(
 activity,
 left_on=["player_id", "first_login"],
 right_on=["player_id", "event_date"],
 how="left"
)

 return merged[["player_id", "device_id"]]

Explanation
groupby + min → finds earliest date.
merge maps each earliest date back to the single row containing the matching device.
Output is trimmed to required columns.

```

## 2.4 Alternative SQL Solutions

### A. Window Function (ROW\_NUMBER)

```

SELECT player_id, device_id
FROM (
 SELECT
 player_id,
 device_id,
 event_date,
 ROW_NUMBER() OVER (
 PARTITION BY player_id ORDER BY event_date
) AS rn
 FROM Activity
) AS t
WHERE rn = 1;

```

### B. FIRST\_VALUE Window Function

```

SELECT DISTINCT
 player_id,
 FIRST_VALUE(device_id) OVER (
 PARTITION BY player_id ORDER BY event_date
) AS device_id
 FROM Activity;

```

### C. Correlated Subquery (less efficient)

```

SELECT
 player_id,
 device_id
 FROM Activity AS a
 WHERE event_date = (
 SELECT MIN(event_date)
 FROM Activity AS b

```

```
WHERE b.player_id = a.player_id
);
```

## 2.5 Alternative Pandas Solutions

### A. sort + drop\_duplicates

```
def game_play_analysis_ll_alt1(activity):
 df = activity.sort_values(["player_id", "event_date"])
 df = df.drop_duplicates("player_id")
 return df[["player_id", "device_id"]]
```

### B. idxmin to retrieve rows directly

```
def game_play_analysis_ll_alt2(activity):
 idx = activity.groupby("player_id")["event_date"].idxmin()
 return activity.loc[idx, ["player_id", "device_id"]]
```

## 2.6 Edge Cases and Discussion

### A. Player logs in only once

Output is trivial: the one device recorded.

### B. Multiple sessions with same event\_date

Impossible because of primary key (player\_id, event\_date).

### C. Ties on earliest date

Not possible.

### D. Large datasets

SQL handles grouping efficiently using indexes.

Pandas may hit memory limits because it loads entire table into RAM.

### E. Time zone considerations

LeetCode always uses plain dates; no UTC/offset concerns.

## 2.7 Complexity Analysis

### SQL Complexity

#### Operation Complexity

---

MIN(event\_date) per player O(N)

Join with main table O(N log N) or better with indexing

Total O(N log N) or O(N)

### Pandas Complexity

#### Operation Complexity

---

groupby + min O(N)

merge O(N)

total O(N)

Memory cost is proportional to the full table size.

## 2.8 Code Review Commentary

### SQL Review Notes

The canonical solution is optimal and widely used.  
Readable and easy to maintain.  
Adding an index on (player\_id, event\_date) greatly improves performance.  
Window function version is clear when more ranking logic is needed.

### Pandas Review Notes

Canonical solution is clean and idiomatic.  
Avoid unnecessary sorting unless using the drop\_duplicates approach.  
Ensure event\_date is a proper datetime type.

### Scalability for Very Large Data

SQL is preferred for data > tens of millions of rows.

Pandas alternatives:

Polars for lazy query-like execution.  
Dask for distributed DataFrame computing.  
DuckDB if SQL is desired within Python.

### 2.9 Example Walkthrough

#### Input

player_id	device_id	event_date
1	2	2016-03-01
1	2	2016-05-02
2	3	2017-06-25
3	1	2016-03-02
3	4	2018-07-03

#### First login dates

player\_id | first\_login

---

1		2016-03-01
2		2017-06-25
3		2016-03-02

#### Match corresponding devices

player\_id | device\_id

---

1		2
2		3
3		1

#### End of Chapter 2

# Chapter 3 — Cumulative Games Played So Far  
\*\*LeetCode 534: Game Play Analysis III\*\*

---

#### ## 3.1 Problem Restatement

We are given the table:

```

```text
+-----+-----+
| Column Name | Type |
+-----+-----+
| player_id | int |
| device_id | int |
| event_date | date |
| games_played | int |
+-----+-----+
PRIMARY KEY: (player_id, event_date)

Each record corresponds to a player's login on a given date, including how many games they played on that day.

Goal:
For each (player_id, event_date), compute:
Total number of games played by that player from their first login up to that date
(i.e., a running cumulative sum ordered by date)

Expected output:
+-----+-----+-----+
| player_id | event_date | games_played_so_far |
+-----+-----+-----+
3.2 SQL Solution (Canonical — Window Function)
SELECT
player_id,
event_date,
SUM(games_played) OVER (
PARTITION BY player_id
ORDER BY event_date
ROWS BETWEEN UNBOUNDED PRECEDING AND CURRENT ROW
) AS games_played_so_far
FROM Activity;

Explanation
The SUM(...) OVER window computes a cumulative running total.
Ordering by event_date ensures correctness.
UNBOUNDED PRECEDING → start from earliest event.
CURRENT ROW → include the current row's games.
This is the standard and optimal SQL formulation.

3.3 Pandas Solution (Canonical)
import pandas as pd

def game_play_analysis_III(activity: pd.DataFrame) -> pd.DataFrame:
# Sort to establish event order
activity_sorted = activity.sort_values(["player_id", "event_date"])

# Compute cumulative games
activity_sorted["games_played_so_far"] = (
activity_sorted
.groupby("player_id")["games_played"]
.cumsum()
)

return activity_sorted[["player_id", "event_date", "games_played_so_far"]]
Explanation
Sorting ensures chronological order.
```

```

groupby(...).cumsum() performs cumulative summation within each player partition.

### 3.4 Alternative SQL Approaches

#### A. Correlated Subquery (inefficient but valid)

SELECT

```
a.player_id,
a.event_date,
(
 SELECT SUM(b.games_played)
 FROM Activity AS b
 WHERE b.player_id = a.player_id
 AND b.event_date <= a.event_date
) AS games_played_so_far
FROM Activity AS a;
```

Complexity:  $O(N^2)$  worst case → not recommended for large tables.

#### B. Self-Join Prefix Sum

```
SELECT
a.player_id,
a.event_date,
SUM(b.games_played) AS games_played_so_far
FROM Activity a
JOIN Activity b
ON a.player_id = b.player_id
AND b.event_date <= a.event_date
GROUP BY a.player_id, a.event_date;
```

Better than correlated query, but still not as efficient as window functions.

### 3.5 Alternative Pandas Approaches

#### A. Using expanding()

```
def game_play_analysis_III_alt1(activity):
 df = activity.sort_values(["player_id", "event_date"])
 df["games_played_so_far"] = (
 df.groupby("player_id")["games_played"]
 .expanding()
 .sum()
 .reset_index(level=0, drop=True)
)
```

return df[["player\_id", "event\_date", "games\_played\_so\_far"]]

#### B. Using .apply (slower)

```
def game_play_analysis_III_alt2(activity):
 df = activity.sort_values(["player_id", "event_date"])
 df["games_played_so_far"] = (
 df.groupby("player_id")["games_played"]
 .apply(lambda s: s.cumsum())
 .reset_index(level=0, drop=True)
)
```

return df[["player\_id", "event\_date", "games\_played\_so\_far"]]

### 3.6 Edge Cases & Subtleties

#### A. Zero games played

games\_played = 0

Still contributes to the cumulative sum (no increment).

#### B. Sparse login days

Players do not log in every day. Missing days do not appear in results.

#### C. Ordering concerns

event\_date must be strictly increasing per player due to primary key constraint.

D. Extremely large datasets

SQL window function is optimized and highly scalable.

Pandas requires entire dataset in RAM → may hit memory limits.

E. Negative values

Not applicable in the LeetCode problem, but cumulative logic still works.

### 3.7 Time & Space Complexity Analysis

SQL

Operation Complexity

---

Windowed SUM O(N) within partitions

Sorting within partitions ~O(N log N) total

Overall O(N log N)

Indexes on (player\_id, event\_date) improve performance significantly.

Pandas

Stage Complexity

---

Sorting O(N log N)

groupby + cumsum O(N)

Total O(N log N)

Memory cost: must hold entire table in memory.

### 3.8 Code Review Commentary

SQL Review Notes

Window-function approach is optimal and idiomatic.

Avoid correlated subqueries in production systems.

Ensure index:

INDEX(player\_id, event\_date)

This accelerates both sorting and partitioning.

Pandas Review Notes

Always sort before cumsum to guarantee correct cumulative order.

Avoid .apply() unless necessary; it is slower and less readable.

Consider using Polars for large-scale cumulative operations:

```
df.sort(["player_id", "event_date"]).with_columns(
 pl.col("games_played").cumsum().over("player_id")
)
```

### 3.9 Example Walkthrough

Input

| player_id | device_id | event_date | games_played |
|-----------|-----------|------------|--------------|
| 1         | 2         | 2016-03-01 | 5            |
| 1         | 2         | 2016-05-02 | 6            |
| 1         | 3         | 2017-06-25 | 1            |
| 3         | 1         | 2016-03-02 | 0            |
| 3         | 4         | 2018-07-03 | 5            |

Processing

Player 1:

2016-03-01 → 5

2016-05-02 → 5 + 6 = 11

2017-06-25 → 5 + 6 + 1 = 12

Player 3:

2016-03-02 → 0

2018-07-03 → 0 + 5 = 5

Final Output

| player_id | event_date | games_played_so_far |
|-----------|------------|---------------------|
| 1         | 2016-03-01 | 5                   |
| 1         | 2016-05-02 | 11                  |
| 1         | 2017-06-25 | 12                  |
| 3         | 2016-03-02 | 0                   |
| 3         | 2018-07-03 | 5                   |

End of Chapter 3

# Chapter 4 — Next-Day Retention After First Login

\*\*LeetCode 550: Game Play Analysis IV\*\*

---

## 4.1 Problem Restatement

We are given the following table:

| Column Name  | Type |
|--------------|------|
| player_id    | int  |
| device_id    | int  |
| event_date   | date |
| games_played | int  |

PRIMARY KEY: (player\_id, event\_date)

Each row represents a player's login activity for a given day.

Goal:

Compute the fraction of players who logged in exactly on the day after their first login date.

Output:

| fraction |
|----------|
|          |

Where:

fraction = retained\_players / total\_players (rounded to 2 decimals)

4.2 SQL Solution (Canonical)

```

WITH first_login AS (
SELECT
player_id,
MIN(event_date) AS first_login
FROM Activity
GROUP BY player_id
),
next_day_login AS (
SELECT DISTINCT
f.player_id
FROM first_login f
JOIN Activity a
ON a.player_id = f.player_id
AND a.event_date = DATE_ADD(f.first_login, INTERVAL 1 DAY)
)
SELECT
ROUND(
COUNT(next_day_login.player_id) / COUNT(first_login.player_id),
2
) AS fraction
FROM first_login
LEFT JOIN next_day_login
ON first_login.player_id = next_day_login.player_id;

```

Explanation

`first_login` → earliest login date for each player.  
`next_day_login` → players who appeared again on `first_login + 1 day`.  
Final query computes fraction with rounding.

#### 4.3 Pandas Solution (Canonical)

```

import pandas as pd

def game_play_analysis_IV(activity: pd.DataFrame) -> pd.DataFrame:
First login date per player
first = (
activity
.groupby("player_id", as_index=False)["event_date"]
.min()
.rename(columns={"event_date": "first_login"})
)

Merge with original table
df = activity.merge(first, on="player_id", how="left")

Compute next day
df["next_day"] = df["first_login"] + pd.Timedelta(days=1)

Identify retained players
retained_players = df.loc[
df["event_date"] == df["next_day"],
"player_id"
].drop_duplicates()

total_players = first["player_id"].nunique()
retained_count = retained_players.nunique()

```

```
fraction = round(retained_count / total_players + 1e-9, 2)
return pd.DataFrame({"fraction": [fraction]})
```

#### Explanation

groupby().min() finds first login.  
Merging aligns each user's later activity with first login date.  
We check if ANY event\_date equals first\_login + 1 day.  
Round to 2 decimals, using small epsilon to avoid float instability.

### 4.4 Alternative SQL Approaches

#### A. Using EXISTS

```
WITH first_login AS (
 SELECT player_id, MIN(event_date) AS first_login
 FROM Activity
 GROUP BY player_id
)
SELECT
 ROUND(
 SUM(
 EXISTS (
 SELECT 1
 FROM Activity a
 WHERE a.player_id = f.player_id
 AND a.event_date = DATE_ADD(f.first_login, INTERVAL 1 DAY)
)
) / COUNT(*),
 2
) AS fraction
 FROM first_login f;
```

#### B. Using Window Functions

```
WITH ordered AS (
 SELECT
 player_id,
 event_date,
 ROW_NUMBER() OVER (PARTITION BY player_id ORDER BY event_date) AS rn
 FROM Activity
),
first AS (
 SELECT player_id, event_date AS first_login
 FROM ordered
 WHERE rn = 1
)
SELECT
 ROUND(
 COUNT(a.player_id) / COUNT(f.player_id),
 2
) AS fraction
 FROM first f
 LEFT JOIN Activity a
 ON a.player_id = f.player_id
 AND a.event_date = DATE_ADD(f.first_login, INTERVAL 1 DAY);
```

## 4.5 Alternative Pandas Approaches

### A. Pivot-like direct membership test

```
def game_play_analysis_IV_alt1(activity):
 first = activity.groupby("player_id")["event_date"].min()
 next_day = first + pd.Timedelta(days=1)

 df = activity.set_index(["player_id", "event_date"])

 retained = [
 p for p in first.index
 if (p, next_day[p]) in df.index
]

 return pd.DataFrame({"fraction": [round(len(retained)/len(first), 2)]})
```

### B. Using merge with adjusted dates

```
def game_play_analysis_IV_alt2(activity):
 first = activity.groupby("player_id")["event_date"].min().reset_index()
 first["next_day"] = first["event_date"] + pd.Timedelta(days=1)

 merged = first.merge(
 activity,
 left_on=["player_id", "next_day"],
 right_on=["player_id", "event_date"],
 how="left"
)

 retained = merged["event_date"].notna().sum()
 total = len(first)

 return pd.DataFrame({"fraction": [round(retained / total, 2)]})
```

## 4.6 Edge Cases & Discussion

### A. Player logs in only once

Cannot be retained; contributes to denominator only.

### B. Player logs in many times

We only care about the day immediately following their first login.

### C. No player logs in the next day

Output is:

fraction = 0.00

### D. All players log in next day

Output is:

fraction = 1.00

### E. Time gaps

If a player logs in on first day, then 5 days later → not retained.

### F. Large datasets

SQL versions scale extremely well.

Pandas version must load entire dataset into memory → may require Dask/Polars for large workloads.

## 4.7 Complexity Analysis

SQL

#### Operation Complexity

---

MIN(event\_date) groupby O(N log N)  
Join / Exists lookup O(N) to O(N log N)  
Overall O(N log N)

Pandas

| Operation   | Complexity       |
|-------------|------------------|
| groupby min | O(N)             |
| merge       | O(N)             |
| filter      | O(N)             |
| Overall     | $^{**}O(N)^{**}$ |

Memory usage: proportional to full table size.

#### 4.8 Code Review Commentary (Optimization Perspective)

##### SQL Review Notes

Canonical two-CTE solution is clean and optimal.  
EXISTS version can be more efficient in some engines.  
Ensure indexing:  
`INDEX(player_id, event_date)`  
Improves join and filtering performance.

##### Pandas Review Notes

Avoid row-by-row loops; use vectorized merge logic.  
Ensure event\_date dtype is datetime for correct arithmetic.  
Use `.drop_duplicates()` to avoid counting players multiple times.

#### Large-Scale Python Processing

Use:  
Polars for high-speed lazy queries  
DuckDB for SQL-on-local-files  
Dask for distributed execution

#### 4.9 Example Walkthrough

##### Input

| player_id | event_date |
|-----------|------------|
| 1         | 2020-01-01 |
| 1         | 2020-01-02 |
| 2         | 2020-05-10 |
| 2         | 2020-05-12 |
| 3         | 2020-08-20 |

First logins

`player_id | first_login`

```

1 | 2020-01-01
2 | 2020-05-10
3 | 2020-08-20
```

Next-day logins

```
player_id | next_day_login
```

```

1 | yes
2 | no
3 | no
```

So:

```
retained_players = 1
total_players = 3
fraction = 0.33 → rounded to 0.33
```

Output

```
+-----+
| fraction |
+-----+
| 0.33 |
+-----+
```

End of Chapter 4